

# Enterprise SW Update

Systems & Software Group (SSG)

# ARM Software Charter



- Lead the Open Source with power efficient solutions
- Investing in community stewardship
- Evolving IP products through Software innovation

Lead



- Timely implementations of ARM-based standards
- Drive open standards which reduce fragmentation
- Producing reference Software

Standardise



- Enable ARM IP
- Add ARM IP support to Virtualisation Software
- Performance & Efficiency optimisation

Enable

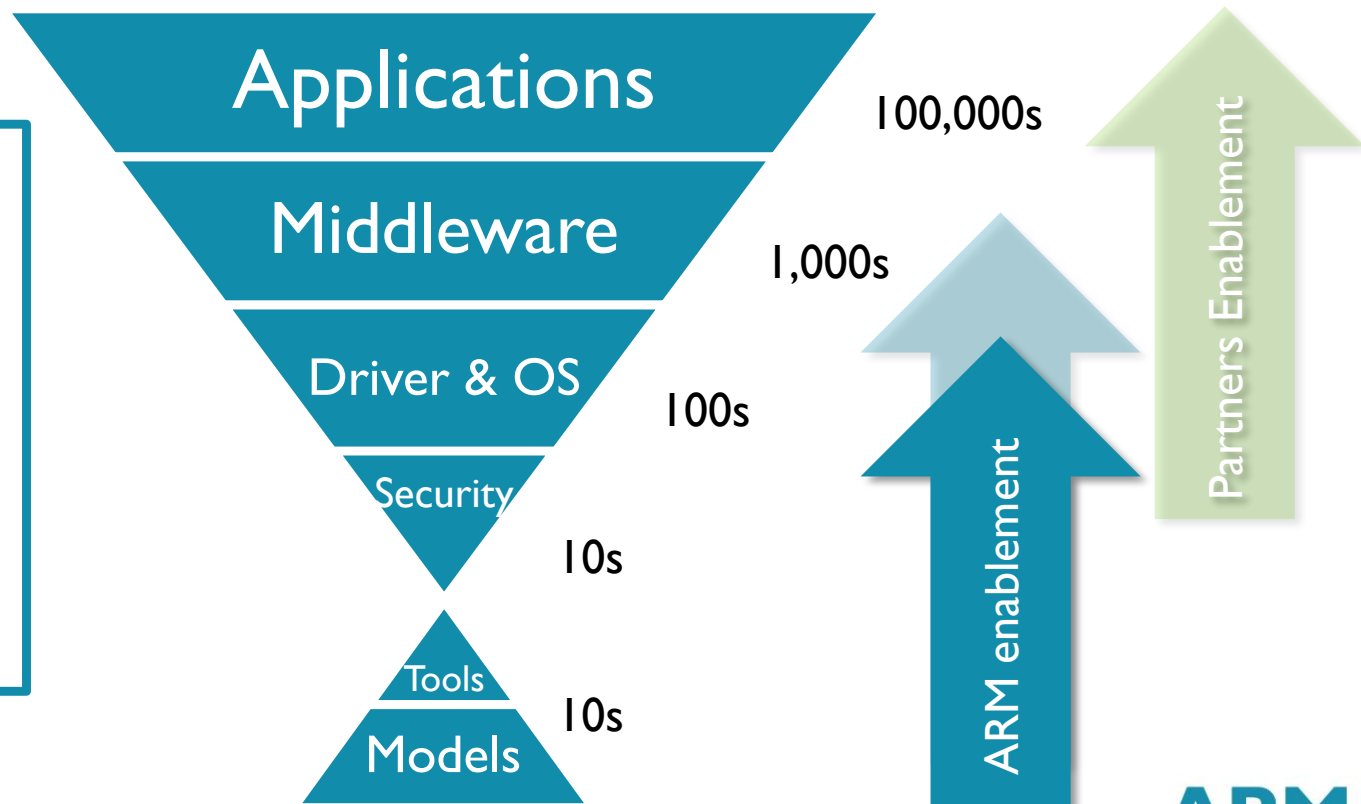


# ARM and Partners Enable the Ecosystem

## ARM Ecosystem and SW Strategy

- A collective effort from ARM and partners
- ARM focuses on enabling key building blocks
- The partnership ensures the entire ecosystem is enabled

▪  is all about this



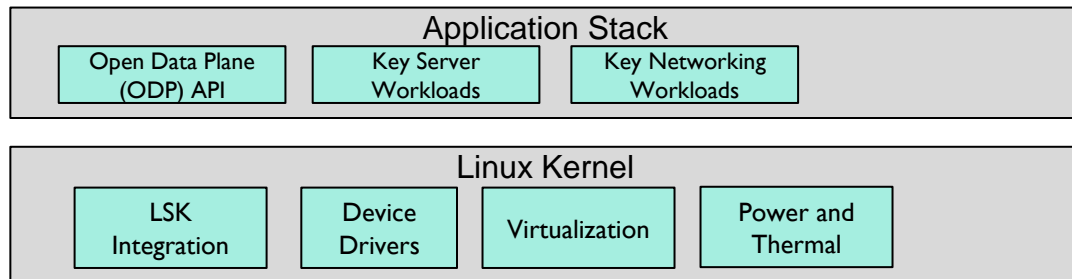
# Subsystems

# Ashbrook Software Overview

## Key

Open Source Software

Closed Source Software



## MCP Firmware

MCP Runtime Firmware

RAS and Manageability

SCP Boot ROM

Platform Boot Initialization

Boot Protocol

## SCP Firmware

SCP Runtime Firmware

System Control & Power Interface (SCPI)

SCP Boot ROM

Platform Boot Initialization

Boot Protocol

## AP Firmware

UEFI

Device Drivers

PI, DXE, RTS

ACPI

U-boot

Device Tree

Uboot-64

Trusted Boot Firmware

Trusted Board Boot

Test Secure-EL1 Payload

Example SMC and Interrupt Handling

AP Boot ROM

Trusted Board Boot

EL3 Runtime Firmware

Secure Monitor Call Handler (SMCCC)

Power State Coordination Interface (PSCI)

World Switch Library

Secure-EL1 Payload Dispatch

# Software deliverables

- System Control Processor (SCP) source, and guide to integration/extension
  - ARM Trusted Firmware source ported to subsystem
  - u-boot or UEFI source ported to subsystem
  - Linaro Stable Kernel (LSK) ported to subsystem
    - Version appropriate to Enterprise – so currently v3.18 LSK based
  - ODP API v1.0 ported to subsystem
  - Release documentation, build instructions
- 
- All with the latest ARM features pre-integrated
    - Security, Power, Performance

# Software Release Details

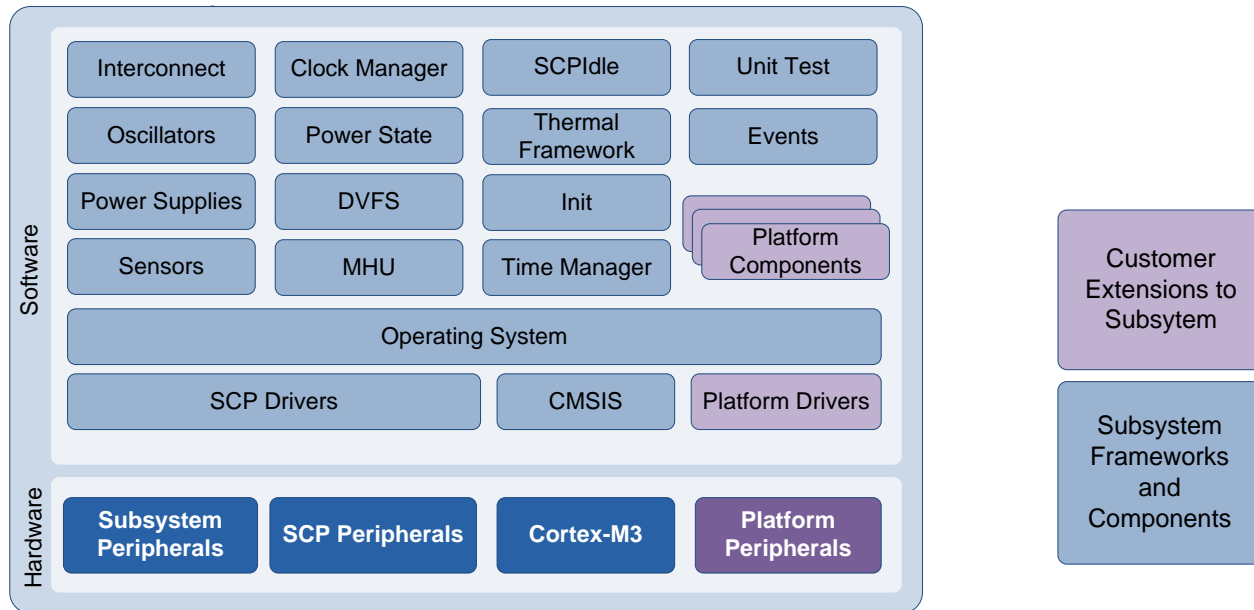
- Subsystem software releases are “development” quality designed to be a starting point for partner software work
  - Integrated all the software common to all SoCs using that subsystem (open-source + proprietary)
  - Validated within ARM's test framework
  - Take the ARM software releases; modify, extend and develop software stack
- For a subsystem we cannot validate all features of a derived SoC
  - Hence all releases are ‘development’ quality
  - Tested on FVP, emulation and ARM Development Platforms (where possible)
  - Test core functionality and ARM provided interfaces
- Maximise use of standard public open-source software
  - Validated on many other systems
  - Our releases contain just the small deltas from the open-source software
  - SCP is not an open-source component, licensed as part of the subsystem

# System Control Processor



# System Control Processor

- SCP Firmware contains
  - Frameworks and components related to the Armstrong design
  - Customer extension components related to the platform built around the Armstrong

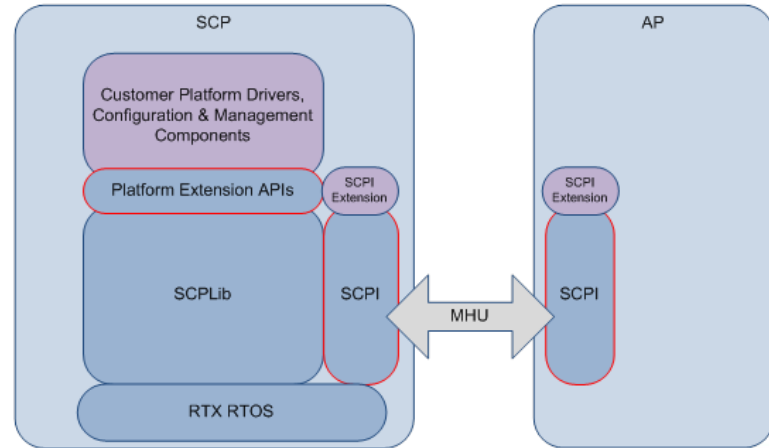


# System Control Processor Overview

- Cortex M3 based subsystem controlled by the AP using the System Control & Power Interface (SCPI) API
- Main responsibilities:
  - Initiating power on sequence and system start up
  - Initial platform configuration
  - Clock and regulator management
  - Servicing power state requests from OSPM
  - Managing a consistent set of hardware states for hardware
  - Handling hardware wake-up requests (timers, interrupts)
  - Managing transitions between operating points to support DVFS
  - Save\restore of state in interconnect and customer integrated DMC during power state transitions
  - Reconciling any capability, performance and wakeup latency requirements with other resource management (PMIC, PLL, sleep on/off/response times)
  - Thermal “loop” (temperature sensor) control

# SCP Interfaces

- Two main interface points to core Library
- AP control of the SCP through the use of the SCPI API and Message Handling Unit (MHU) driver, with the ability to add customer's platform-specific commands
- Platform specific extensions via the Platform Extension APIs to allow integration of platform-specific components
- Platform Drivers for customer-specific hardware under SCP control
  - e.g. Power Supplies, Oscillators, Sensors
- Platform Management Components
  - Abstract platform drivers from SCP components
  - Enables SCPI extension handling
- Configuration
  - Start-up options
  - Memory layout



# Message Handling Unit (MHU)

- Allows inter-processor communication between SCP and AP
- 31-bit register area for messages + interrupt mechanism
  - For messages that require extra payload, (Non-Secure) Scratch RAM and Secure RAM are used
- 2 priority levels (High and Low) plus a secure channel in both directions
  - Up to 6 simultaneous messages
  - Secure messages processed first, then high priority, then low priority
  - Secure channel only accessible to secure world software
  - Independent to channel priority (in the software)
- Messages are decoded into SCPI commands

# Manageability Control Processor (MCP)

# What is an MCP?

- MCP stands for Management Control Processor
- The MCP is separate IP block that provides the following functionality:
  - On chip management
  - Agentless management
  - RAS on chip
  - Separation of trust domains (MCP = non-trusted; SCP = trusted)
  - Provides single interface to the platform BMC
- MCP current implementation is a Cortex M7

# Why incorporate MCP into SOC design?

- Create Enterprise Standards

- BMC to SOC communications are not standard today. The MCP provides a means to standardize this interface by reducing the need for BMC and platform vendors to support different interfaces for every SOC vendor.

- ARM Provided Reference Code

- ARM will provide reference code with the MCP to guide partners on how to implement on-chip management and RAS
- SOC vendors can choose to use reference code as a base for their own implementation

- Prepare for the Future

- BMC operations will eventually be consolidated into the SOC as having a separate BMC on the platform increases the platform cost. The MCP provides a place for BMC vendors to place their code as this transition takes place.
- The Cortex-M7 has an option for a GbE interface that can leverage a network accelerators NCSI interface for out-of-band communications. For low cost servers, this could remove the need for a platform BMC.

# RAS/Manageability



# RAS in ARM Architecture

- ARM architecture for A processors has had very little to say about RAS functionality prior to v8.2
  - Errors reported to software via fault status register and aborts and system error interrupts (SEI)
  - Additional important information on the faults are in implementation-specific registers in the core and in devices (such as memory controllers)
- But, many implementations had RAS features
  - Primarily, ECC on caches and main memory
- Additional RAS features could have been implemented as well, such as
  - Memory scrubbing
  - Bus error protection
  - Fault isolation
- One useful RAS mechanism hasn't usually been implemented, an error fence to aid in containing errors to the software that caused them

# v8.2 RAS Extensions Summary

## A language for describing RAS

- Detection
- Consumption
- Containment
- Isolation
- Recovery
- Poisoning

## Error Reporting

- Synchronous and asynchronous external aborts (SError)
- Error recovery interrupts
- Firmware-first error handling
- Standard syndrome registers for error handler

- Provides a framework for building RAS features
- Few actual additions to architecture
- Does not specify other RAS features in implementations

## Error Memory Barrier

- For isolating uncontainable errors to the software environment that caused them.
- EMB – new error barrier instruction
- Issued before switching VM or process

## Fault Handling

- Fault handling interrupts
- Error records
- *Today's systems require implementation-specific software*

- Required for v8.2
- Status: v8.2 is beta

# RAS Architecture Extensions

- Give tools to build RAS enabled systems, but implementations need to fill in quite a bit of detail:
  - Which portions of the design should be protected and against which errors?
  - How to join IP-level RAS features with system-level behaviors?
  - Where to send fault handling and error recovery interrupts?
  - Is it possible to recover from a particular error and how to do so?
  - Etc.
- So, quite a bit is needed to bridge the low-level tools of the RAS extensions to the needs of an end system

# Hardware for RAS

- Error detection
  - Detectors as appropriate for system design and for FIT goal
  - Cover most likely fault sources (e.g. DRAM data corruption or watchdog timeouts)
- Error correction
  - Some detected errors can be corrected; and correction lowers FIT rate
  - Correction can be performed by hardware, firmware or software (most usually hardware)
- Signaling software – Faults and errors are signaled to software via:
  - Abort
  - System error interrupt
  - Fault handling interrupt<sup>†</sup>,
  - Error recovery interrupt<sup>†</sup>

<sup>†</sup> New in the v8.2 RAS extensions

# Hardware for RAS (2)

- Recording information about the error
  - At the point of detection, usually store information about the error into registers
  - RAS architecture extensions give standard ways to access this information and slightly standardize the information recorded
- Signaling deferred errors
  - A system that supports deferred errors will need a way to signal poison between IP
- Logging
  - Errors need to be logged in non-volatile storage, so such storage must exist (somewhere)

# Legacy vs v8.2

- In v8.2, when RAS extensions are in the architecture, RAS becomes more standard
  - Defined interfaces for signaling errors to software and for getting error reports
  - New instruction ESB to fence errors before changing software environment (e.g.VM switch)
- Everything RAS could be done by implementation-specific means today

# Upcoming RAS hardware features

- ARM IP will implement the following features for LAC/EAC products in 4Q 2015 and 1H 2016:
  - Interconnect that implements data error protection throughout – byte parity on transport and ECC in the system cache
  - Data poisoning on 64-bit data – interconnect, memory controller, caches, CPUs
  - Memory controller Reed-Solomon advanced error detection and correction per cache line that can correct 4-bit memory chip failures
  - ARM v8.2 RAS error reporting – CPUs, interconnect, memory controller, GIC

# PCIe RAS

- PCIe is important to the system reliability and serviceability
- ARM does not offer PCIe root port IP
- The PCIe root port should be selected for the desired RAS features

We believe that the following PCIe RAS features are desirable for a server product to be competitive:

- Advanced error reporting (AER)
- End-to-end CRC (ECRC)
- Downstream port containment and enhanced DPC (DPC and eDPC)
- Hot plug



# RAS in Subsystems

# System-Level Standards

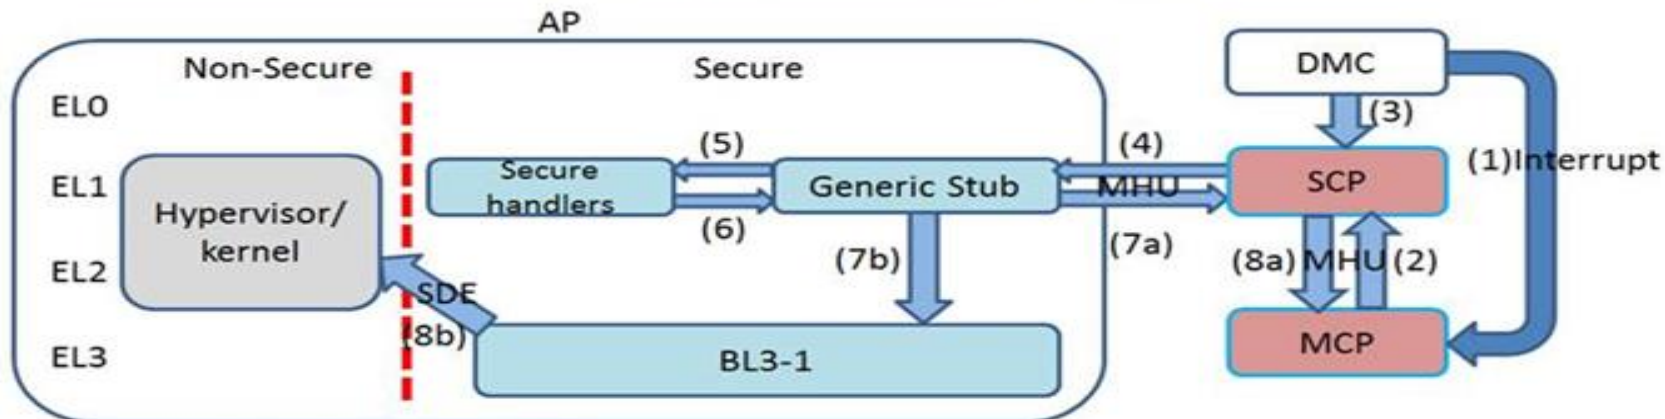
- IPMI 2.0/DCMI 1.5 reference code
  - Power on/off/status
  - FRU Data
  - System Event Log (SEL)
  - SOC Power/Thermal Data to BMC
  - In-band communications messages via SCP Mailbox (MHU)
  - Future out-of-band interface standard for BMC<->MCP communications in development
- Alert events according to ASF v2.0
- Firmware update according to PIC HPM.1
- SBSA Level 3
- ARM RAS HW Standards for v8.0 and later

# ARM Subsystems RAS support

- ARM subsystems will use ARM IP components supporting RAS features from previous slides
  - Ashbrook is first subsystem with many of the RAS features
- In addition, will have firmware support for RAS that coordinates error logging and handling between main processors, embedded system controllers (SCP and MCP) and an external BMC.
  - Primary function of SCP is power control
  - Primary function of MCP is management interface
  - Primary function of BMC is platform management
  - RAS firmware on the main processors (application processors) handles RAS interrupts and controls software recovery actions
- See the S/W slides next.

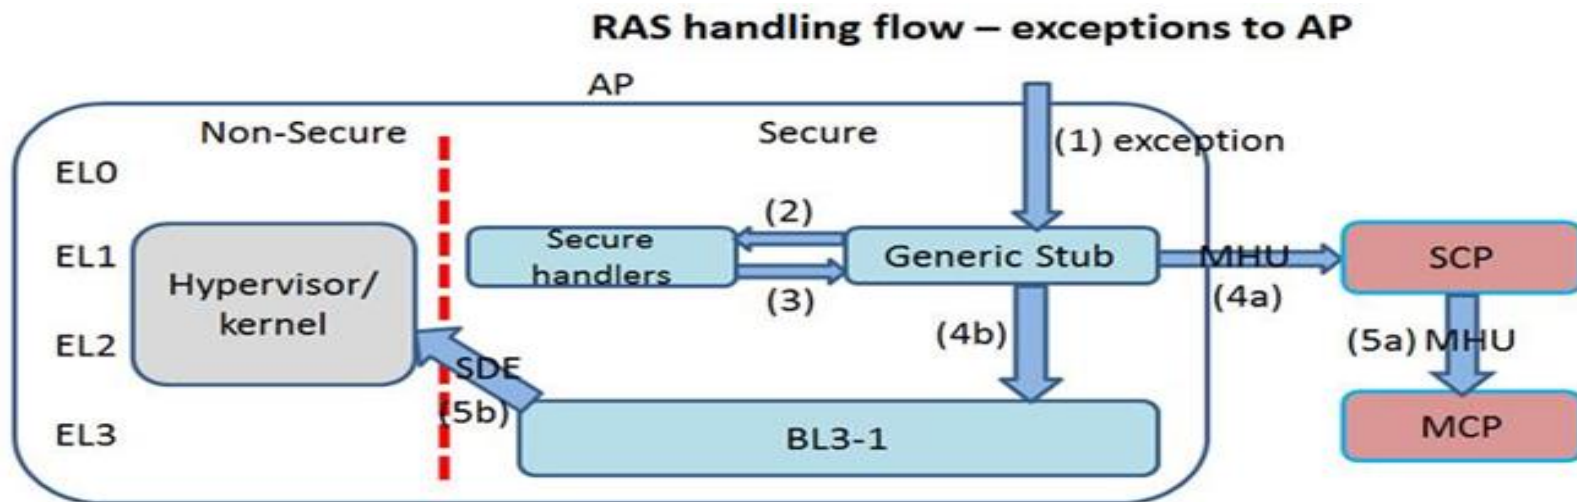
# MCP-Directed RAS Flow

RAS handling flow – exceptions to MCP



- (1) DMC interrupt comes to MCP
- (2) MCP requests SCP to handle DMC event
- (3) SCP reads relative DMC registers
- (4) SCP reports DMC event to Generic Stub with DMC registers
- (5) Generic Stub calls secure handlers (chain)
- (6) Secure handlers create the CPER for this error and return back to Generic Stub  
According to policy from handler, Generic stub may route the CPER to :
  - (7a) SCP, then SCP sends the CPER to MCP (8a)
  - (7b) Hypervisor/kernel using SDE interface (8b)

# ARM TF-Directed RAS Flow



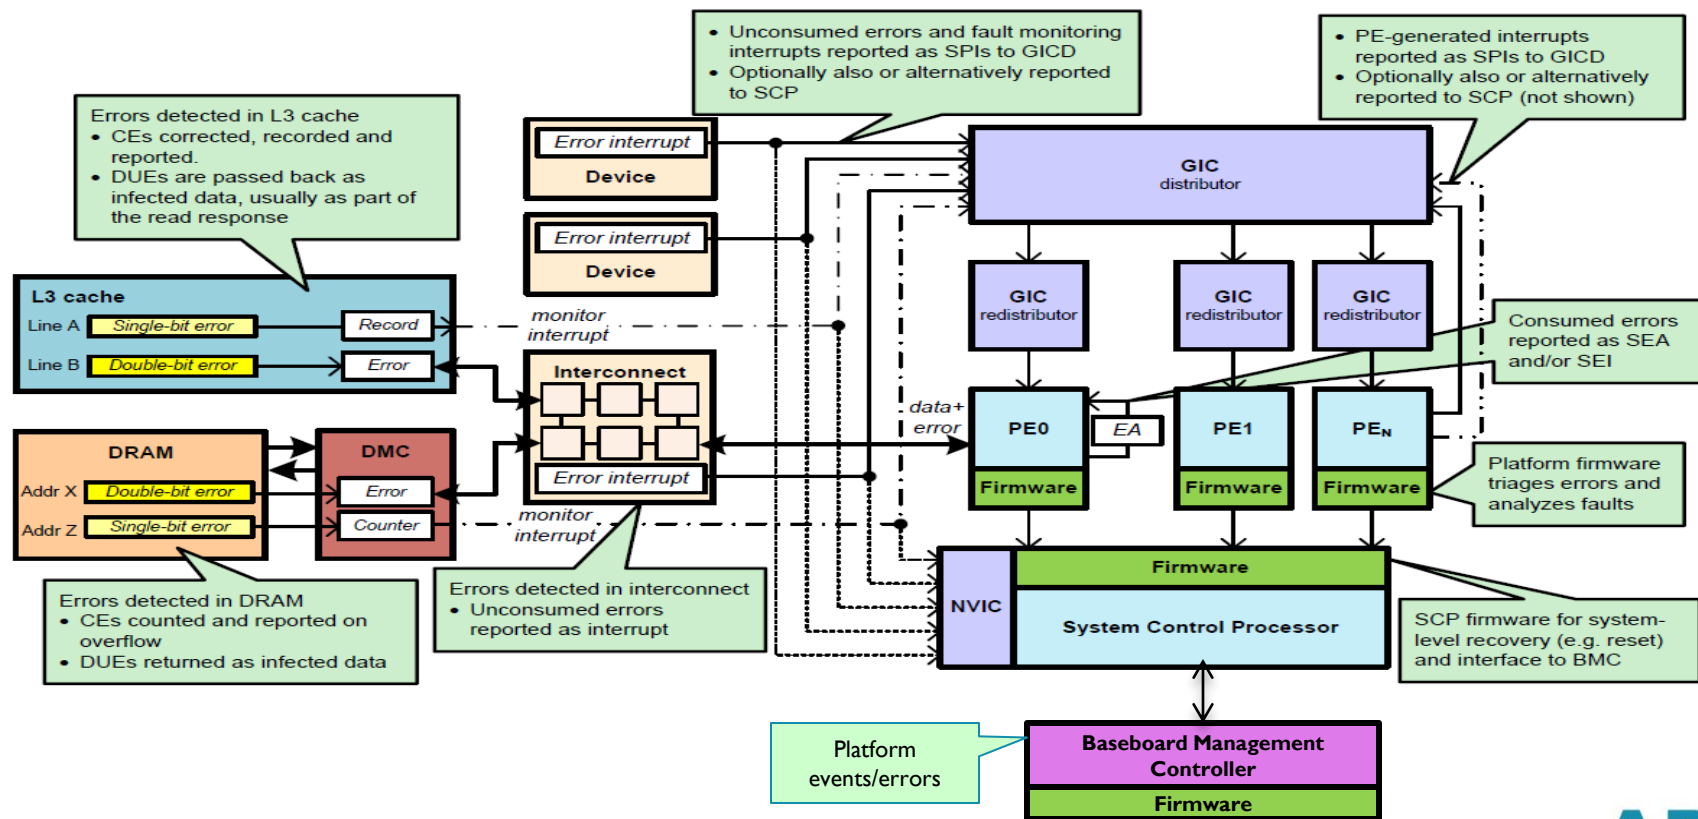
- (1) Exception comes to AP (Generic Stub)
- (2) Generic Stub calls secure handlers (chain)
- (3) Secure handlers create the CPER for this error and return back to Generic Stub

According to policy from handler, Generic stub may route the CPER to :

- (4a) SCP, then SCP sends the CPER to MCP (5a)
- (4b) Hypervisor/kernel using SDE interface (5b)

# Software RAS

# System View of RAS



# SW RAS – What ARM is working on

- Working with partners to build a Minimum ARM Enterprise SW RAS Standard that will complement the ARM HW RAS spec that is in review
- Delegated exceptions for communicating asynchronous error events to OS and hypervisor
- Reviewing proposals for change to standards and SW (ACPI, UEFI, ARM Trusted FW, SBBR, hypervisors, perf) and providing feedback
- Working with Linaro to review and update LEG cards regarding RAS

## NOTE

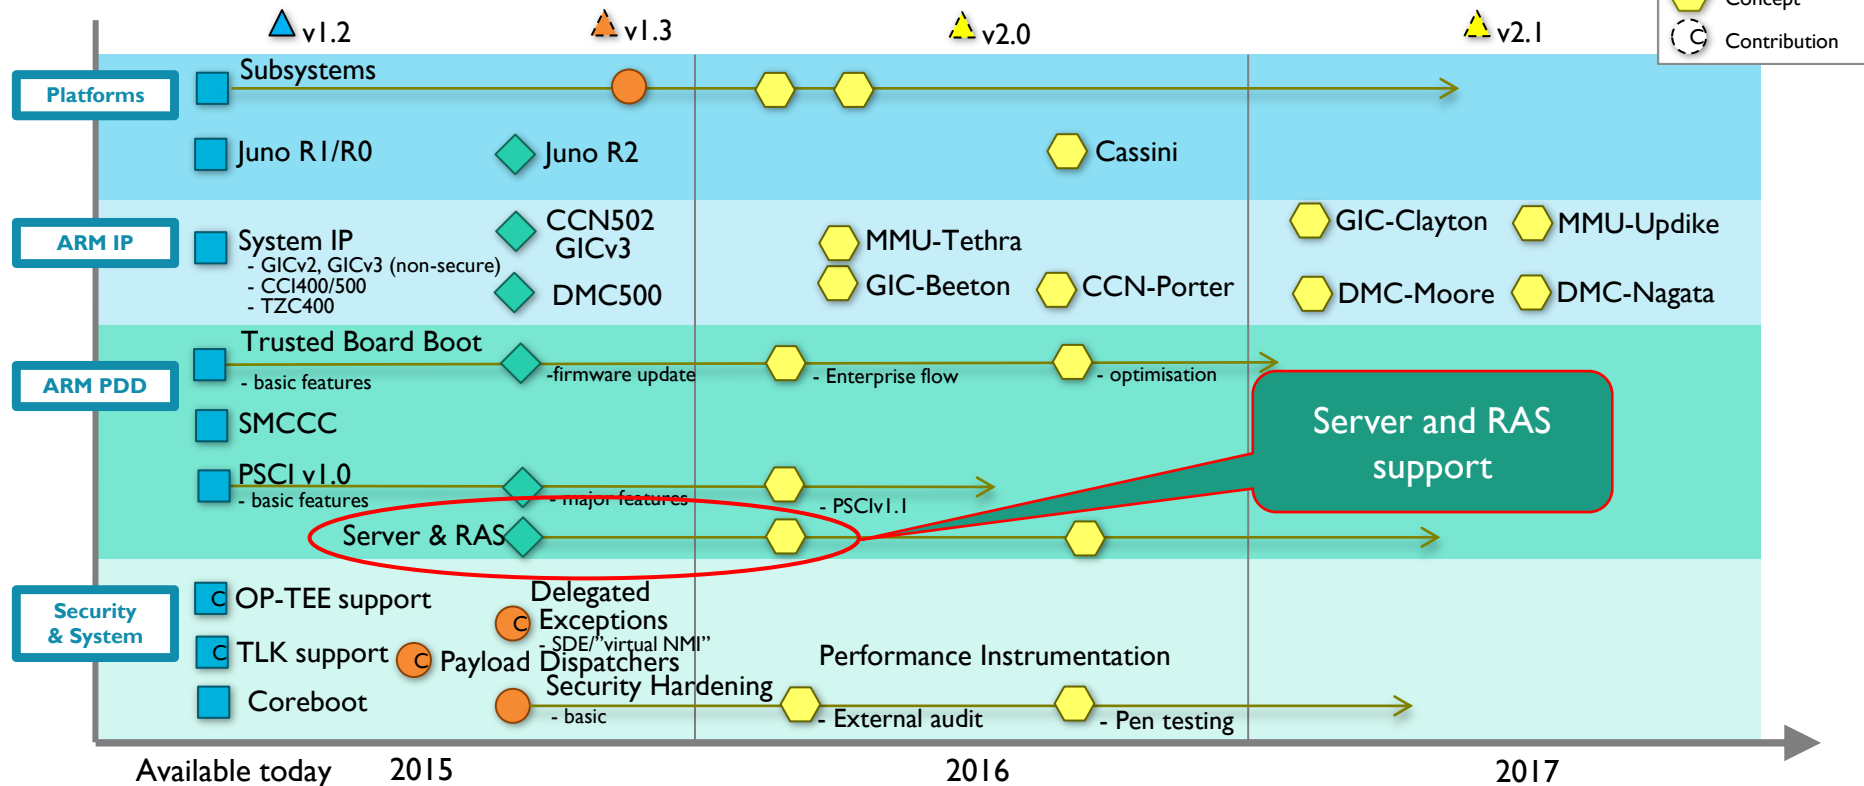
ARM also plans to work with one or more lead partners to ensure the ‘Minimum SW RAS Standard’ has an open source implementation available tied to physical hardware



# ARM Trusted Firmware Support for RAS

- ARM Trusted Firmware (TF) support for RAS:
  - Receive fault reporting and error recovery interrupts
  - Collect information into common platform error record (CPER)
  - Route to handlers in TF, silicon vendor and OEM code
  - Send logging information to BMC
- First TF implementations in 2016

# ARM Trusted Firmware\* Roadmap



\* <http://github.com/ARM-software/arm-trusted-firmware>

# What Work is Needed

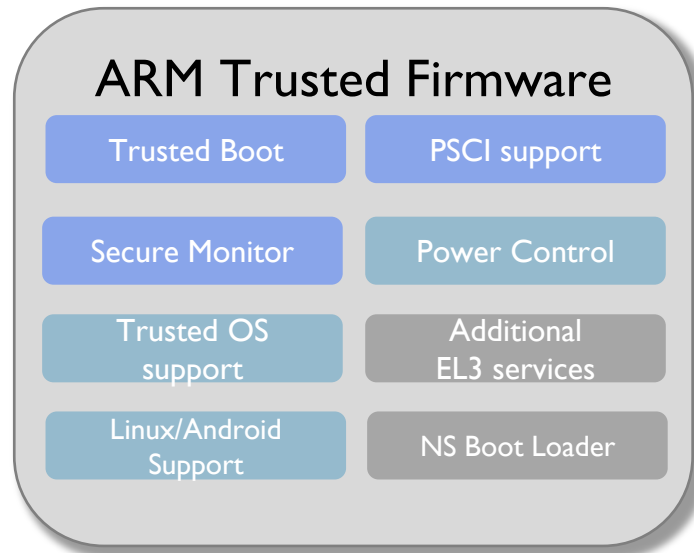
- Standard Error Reporting Mechanisms
  - UEFI Common Platform Error Record (CPER) needs to be updated
  - UEFI Run-time services support defined
  - ACPI Platform Error Interface (APEI) needs to be updated
  - Synchronous Error Aborts (SEA) for contained consumed errors
  - Asynchronous external aborts (AEA) for (uncontained consumed errors)
    - Needs to report the level of contamination
  - Interrupts for all other errors need to be defined
  - Common Record format for RC/EL3 to BMC fault reporting
- Reference implementation of SDE (Software Delegated Exception) in EL3 FW
- Routing of RAS Interrupts
  - Initial Routing of interrupts to EL3 where a small amount of monitor code will redirect the interrupt to either Secure EL1 or EL2/EL1 for processing
- Controls for Routing Synchronous External Aborts to EL2
  - Today's systems require EL1 code to pass a synchronous external abort up to FW, which isn't truly 'firmware first'
- Secure Virtual SEI for EL3 to send an SEI to lower EL's
  - High Priority interrupt needed
  - Registers to define ESR on taking Secure or Non-secure virtual SEIs
  - Secure EL1 Partitions?
- Changes to Perf for error/fault management within OS
- Hypervisor fault handling defined (Type 1 and Type 2 hypervisors have different requirements)

# Backup

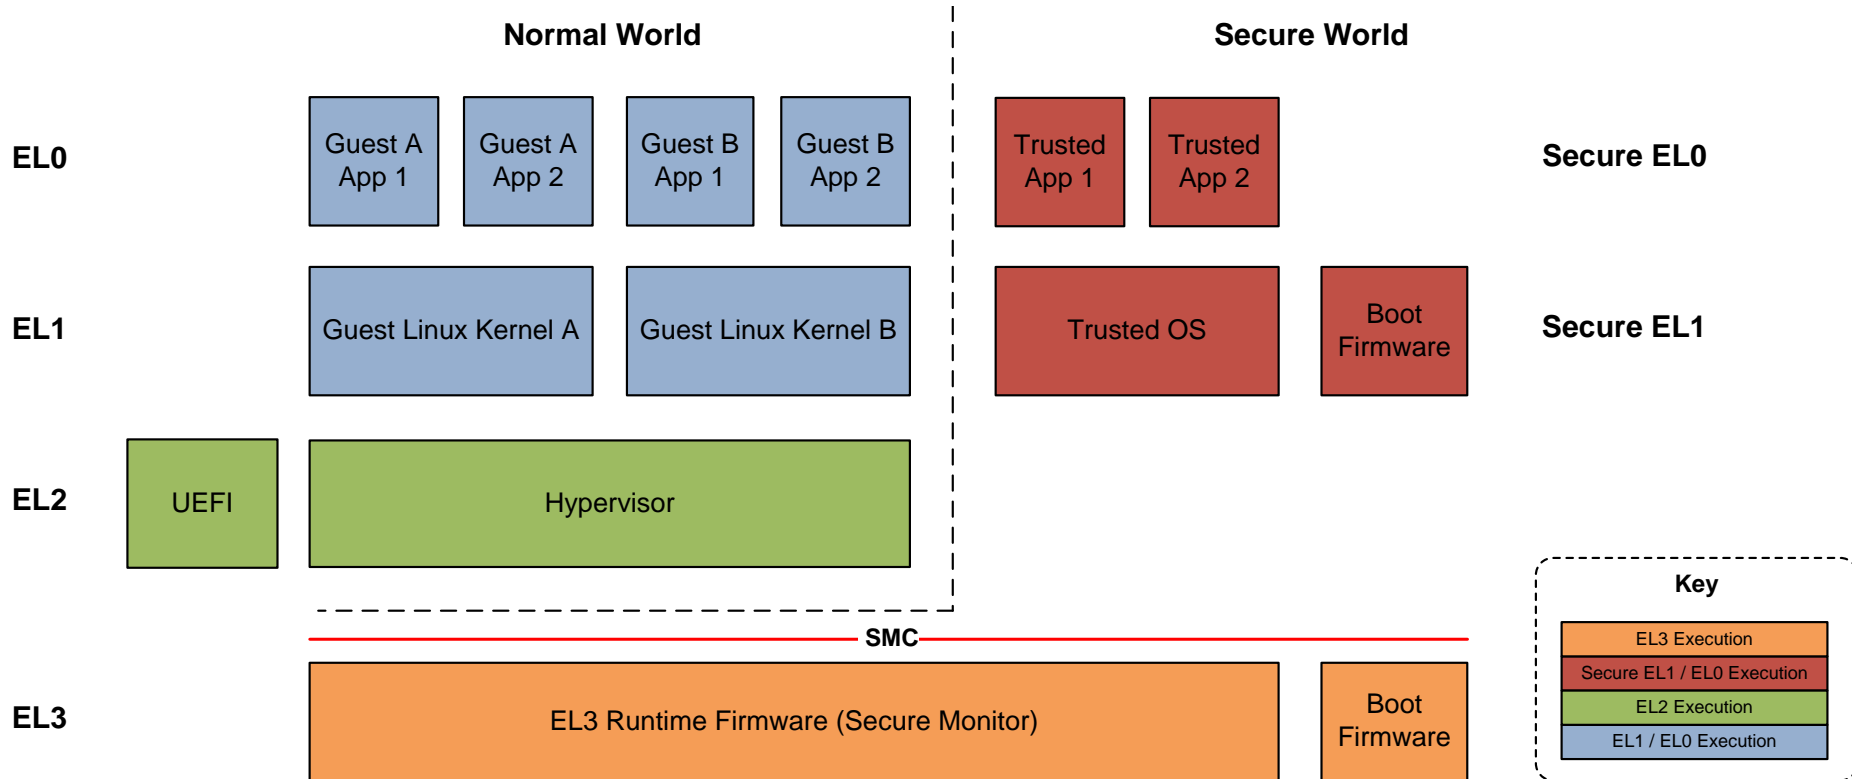
# ARM Trusted Firmware

# ARM Trusted Firmware

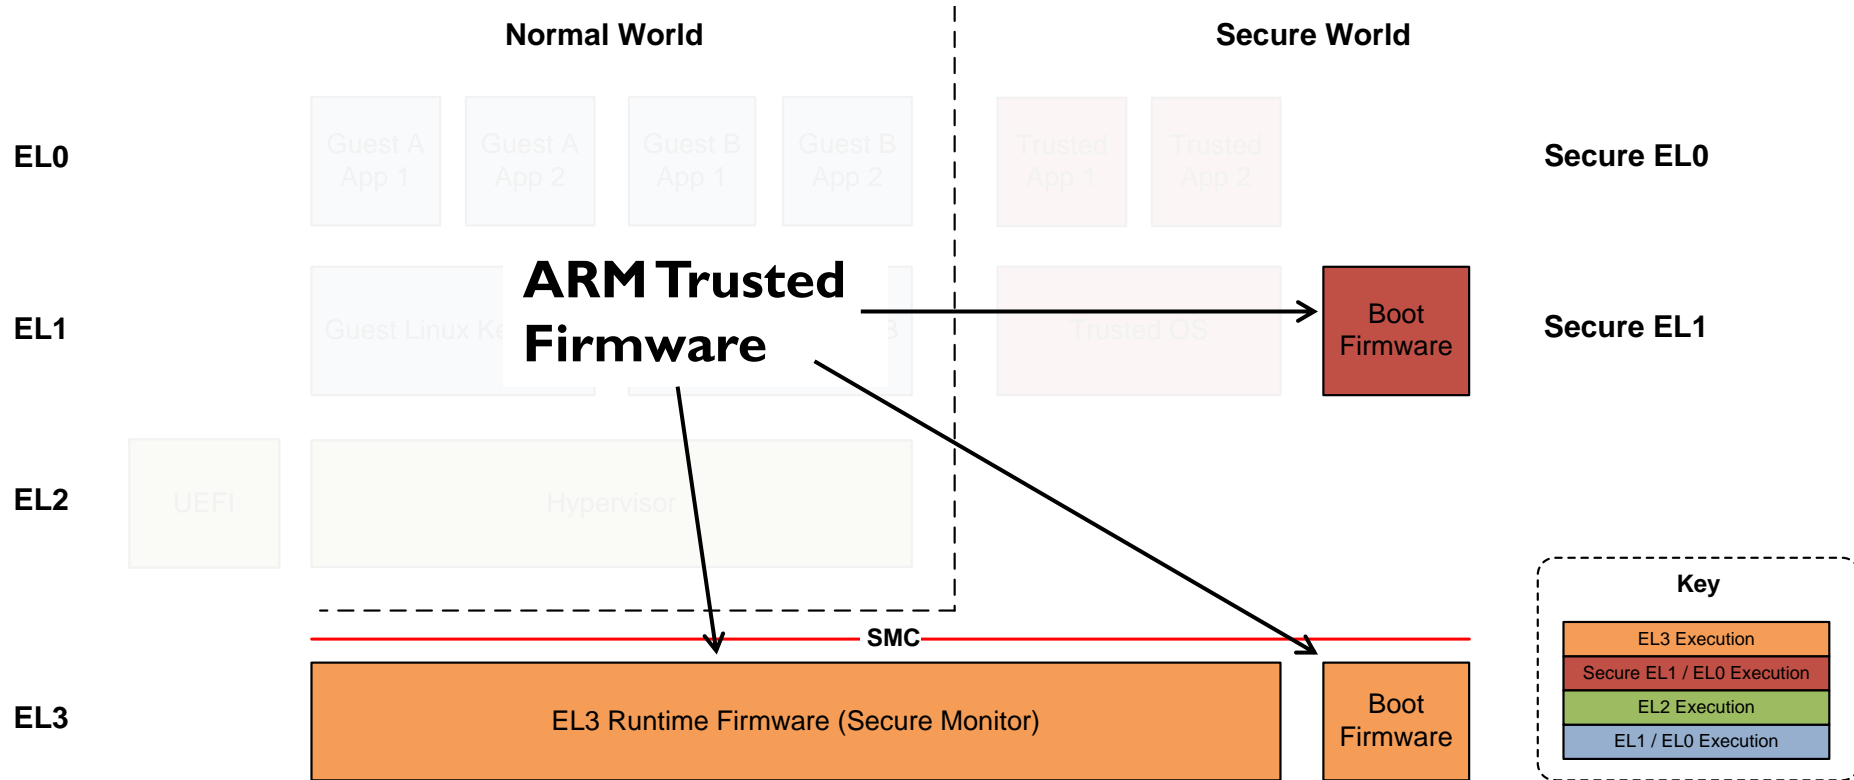
- 64bit Boot and Runtime Firmware for ARMv8-A
  - Mostly executes within the EL3 exception level
- Trusted boot sequence
- Runtime CPU power interfaces (PSCI), used by Linux
- Flexible Secure OS integration
- Extensible framework for SiP and OEM
- Open Source Project with BSD license and standard CLA
  - Managed by ARM, hosted at GitHub
  - <https://github.com/ARM-software/arm-trusted-firmware>
  - <http://www.arm.com/community/open-source-contributing.php>



# What is ARM Trusted Firmware?

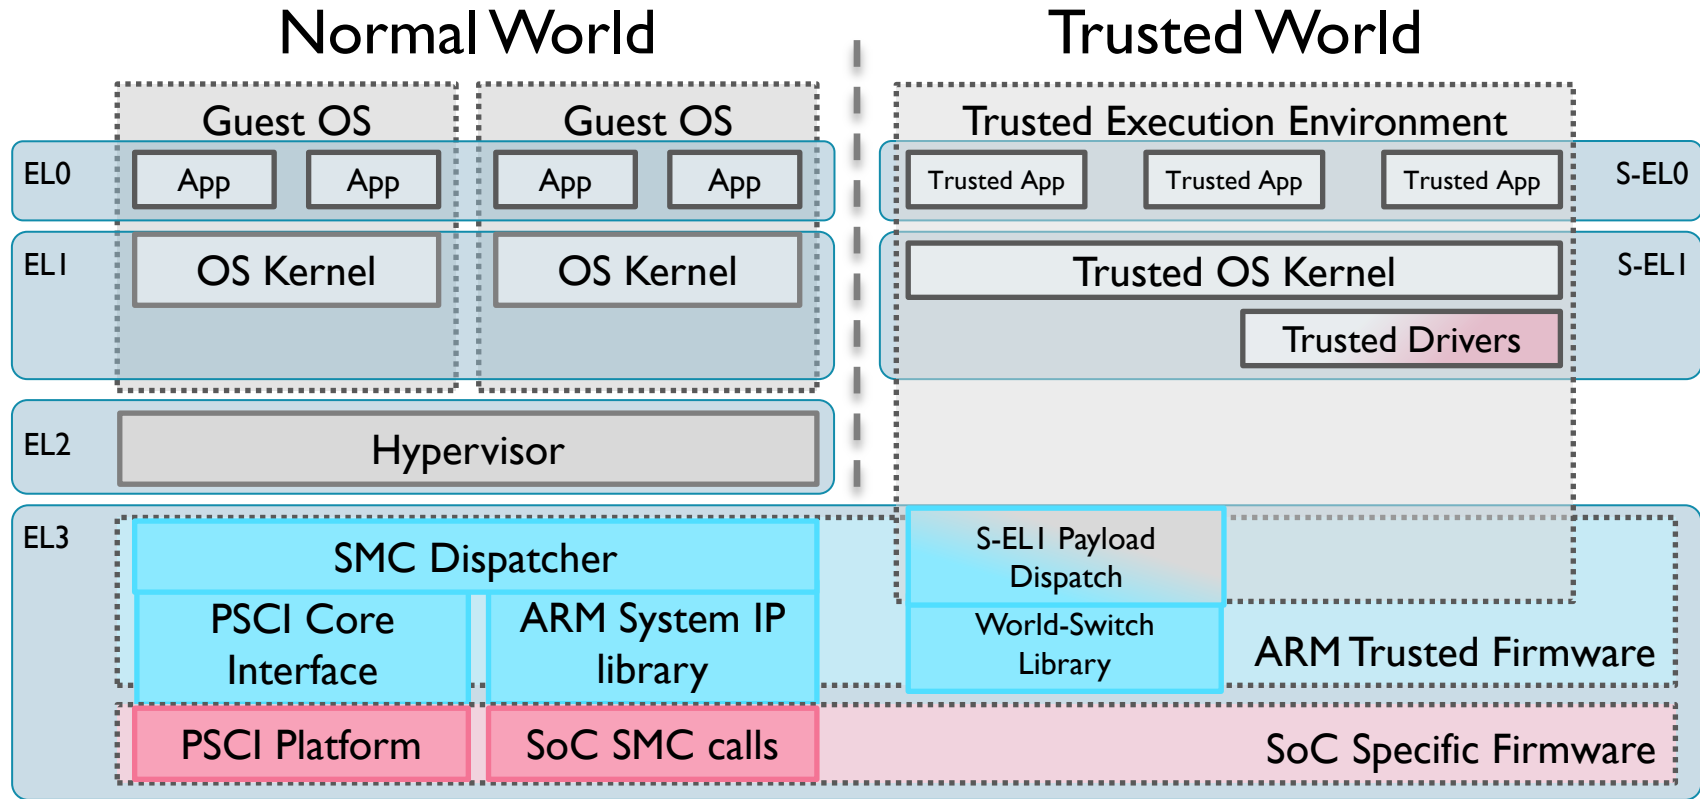


# What is ARM Trusted Firmware?





# ARM Trusted Firmware Architectural Overview

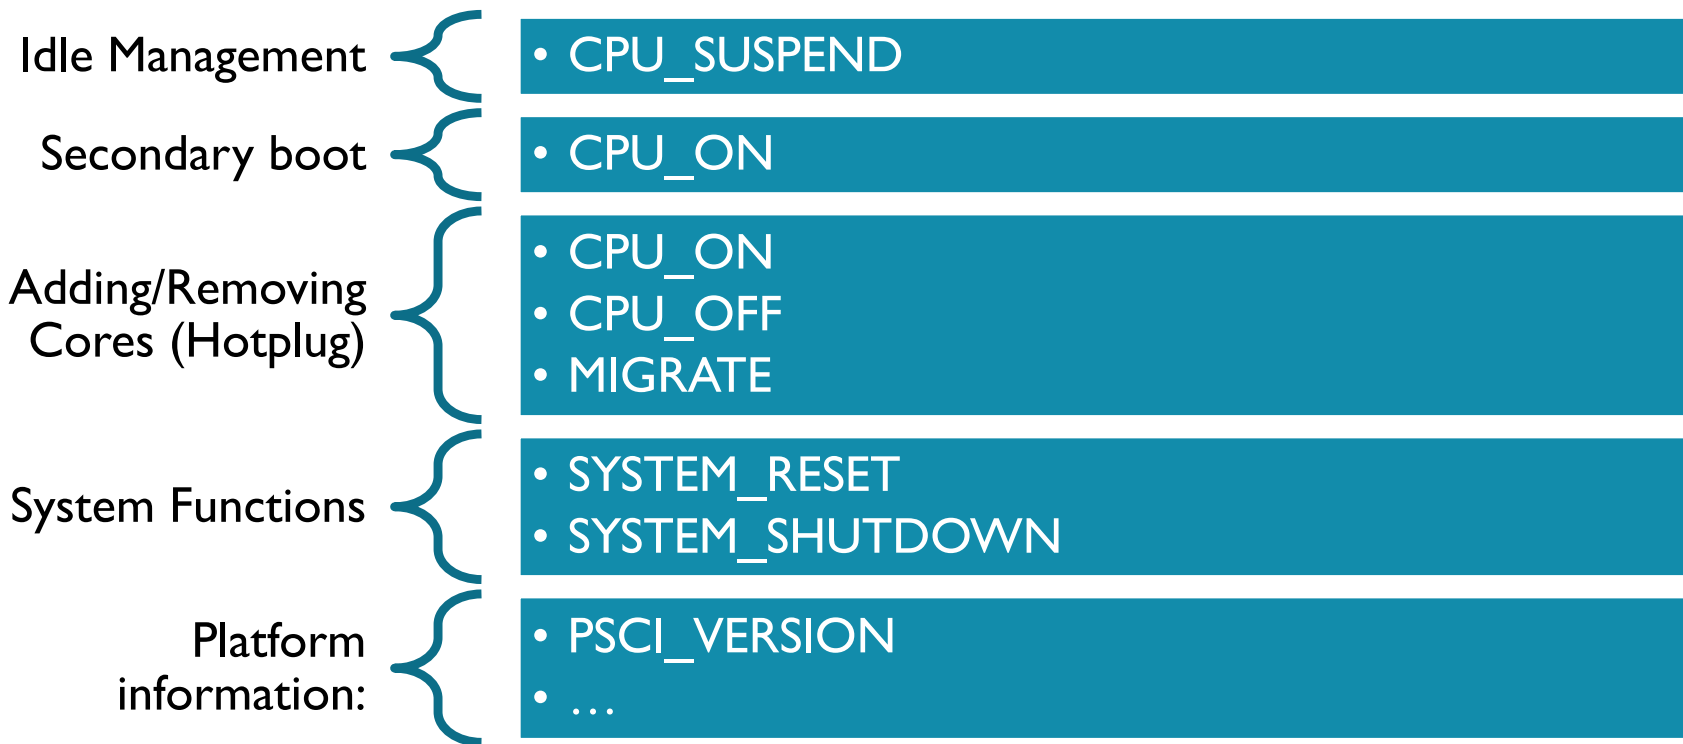


# Power State Coordination Interface (PSCI)

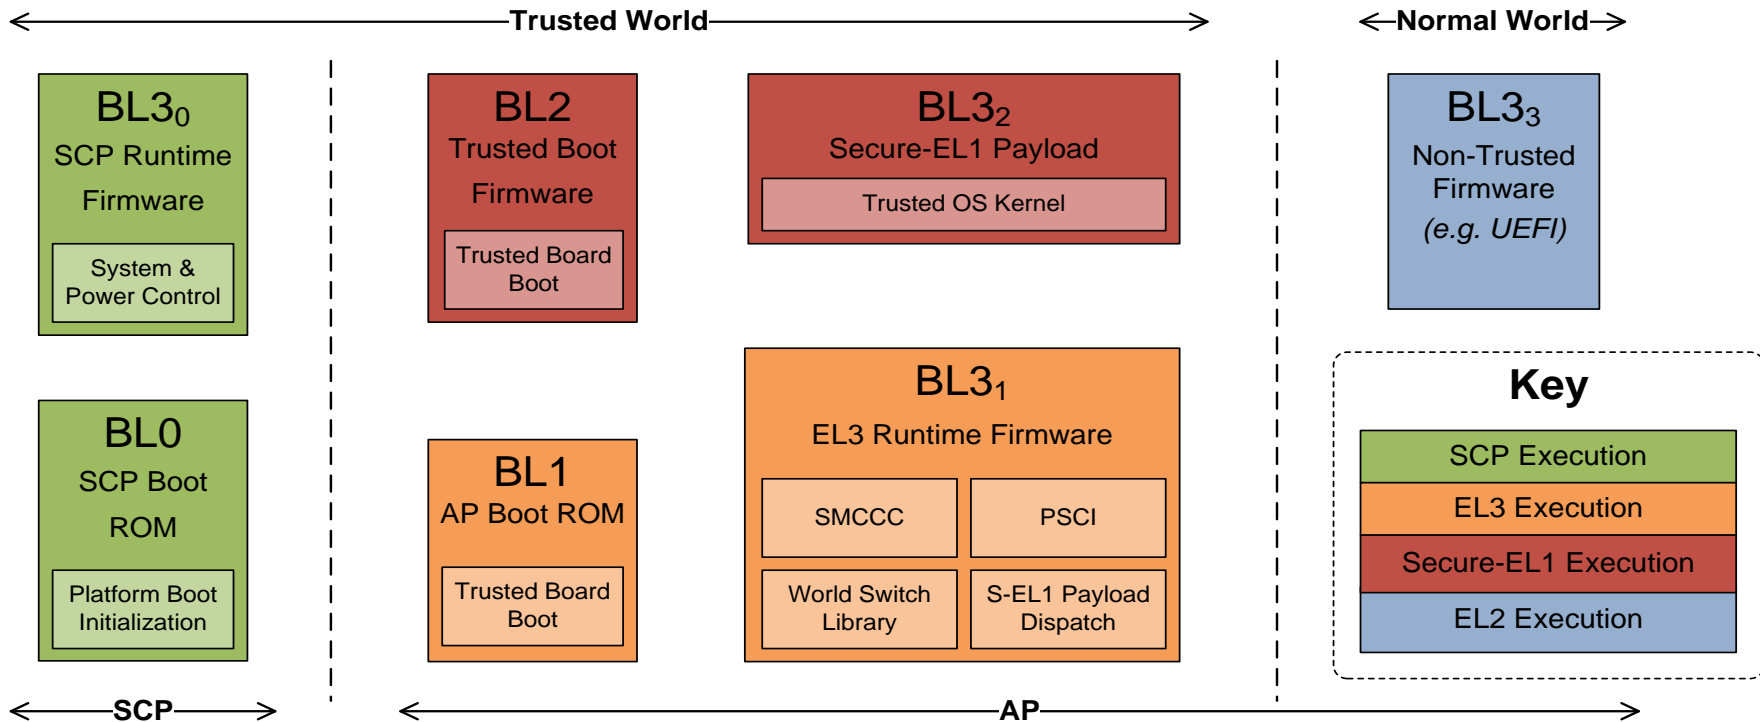
- ARM standardised interface for OS power management
- Supports virtualisation and power management of Trusted OS
- Main method for power control in Linux AArch64 kernel
- Draft specification (v0.2) available today in the ARM Information Center:  
<http://infocenter.arm.com/help/topic/com.arm.doc.den0022b/index.html>
- v1.0 planned for Q4



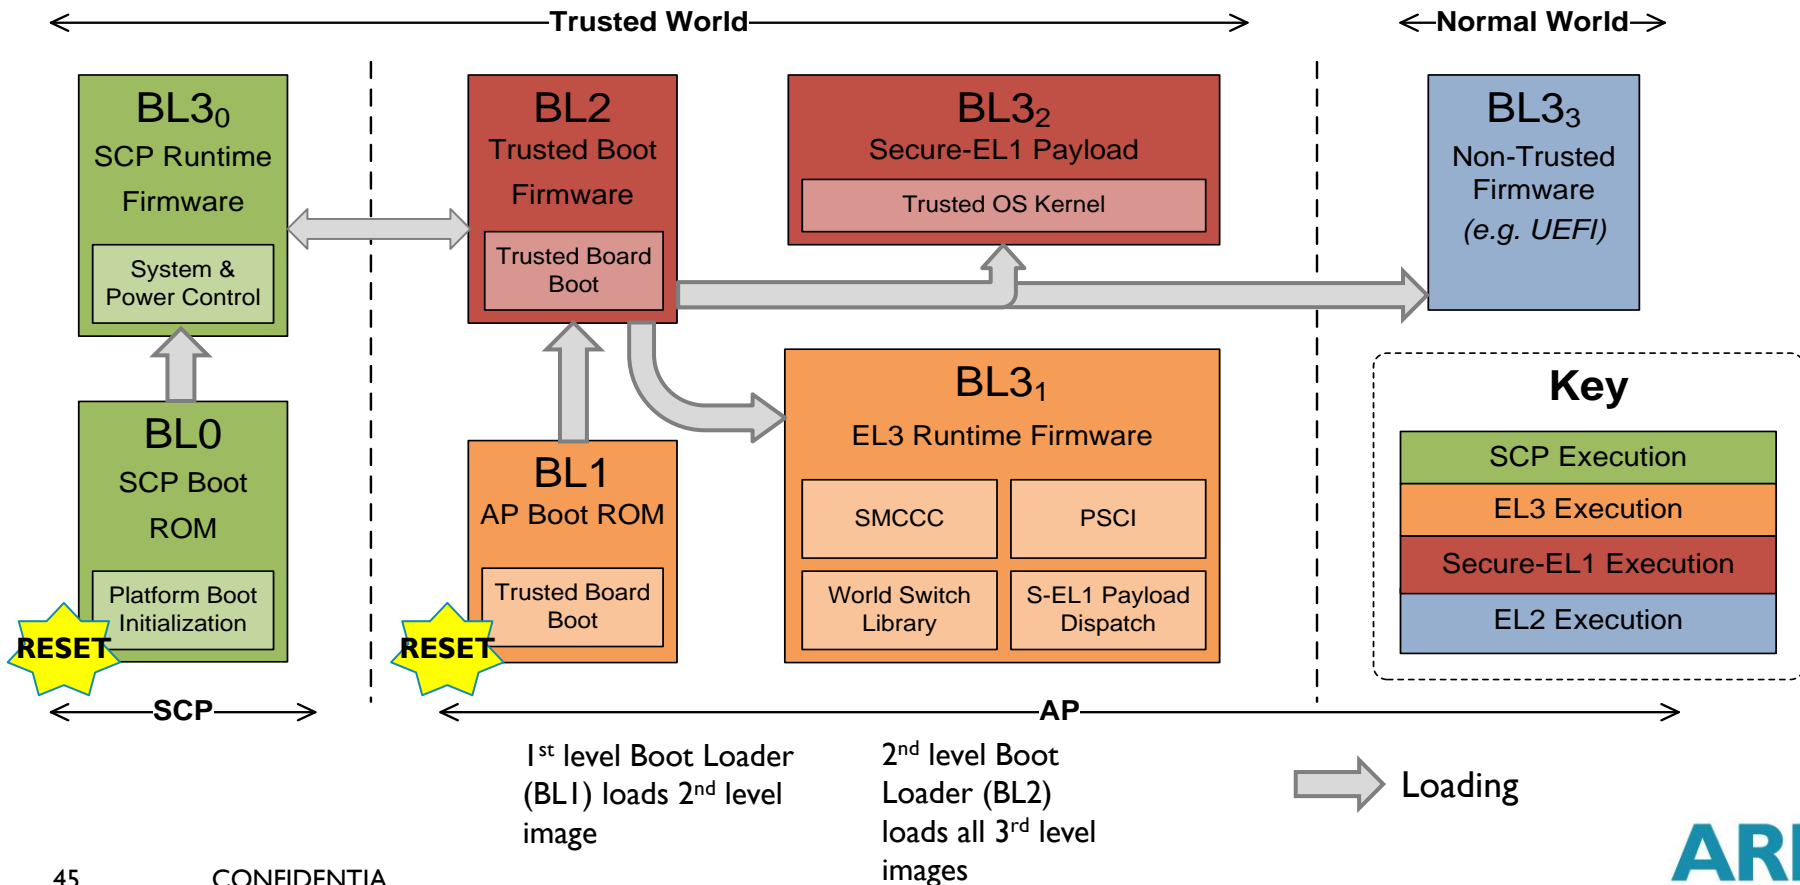
# PCSI Functions



# Trusted Boot Image Terminology



# Trusted Boot Image Loading Flow

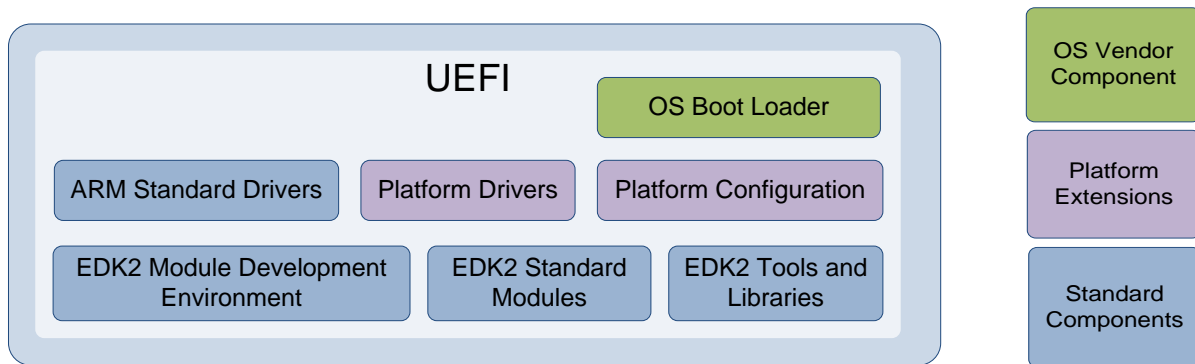


# Firmware Image Packages (FIPs)

- All non-ROM firmware images stored as a FIP in Non-Volatile Memory (NVM)
  - For now: BL2, BL3<sub>0</sub>, BL3<sub>1</sub>, BL3<sub>2</sub> and BL3<sub>3</sub>
  - Later: Recovery firmware images and certificates (for TBB)
- FIP begins with a Table of Contents (ToC)
  - A header plus an entry per image/certificate
- BL1 and BL2 read FIP contents to load images/certificates into Trusted RAM
- A tool (and makefile support) is provided for FIP creation/enquiry

# UEFI

- UEFI specs owned by UEFI forum
  - AArch32 supported since v2.3. AArch64 supported since v2.4.
- EDK2 is open source implementation of UEFI forum specs
  - Large developer community <https://github.com/tianocore/edk2.git>
  - ARM maintains the ARM architecture and contributes to the overall infrastructure. Linaro maintains a platforms tree and makes associated monthly releases [to the tree]
- Modular nature of EDK2 makes it easy to re-use and customize for each platform
  - Minimizes platform specific porting



# SW Workloads



# ARMv8-A Infrastructure Ecosystem Building Momentum

## Key Applications



## Middleware



## Operating System & Virtualization



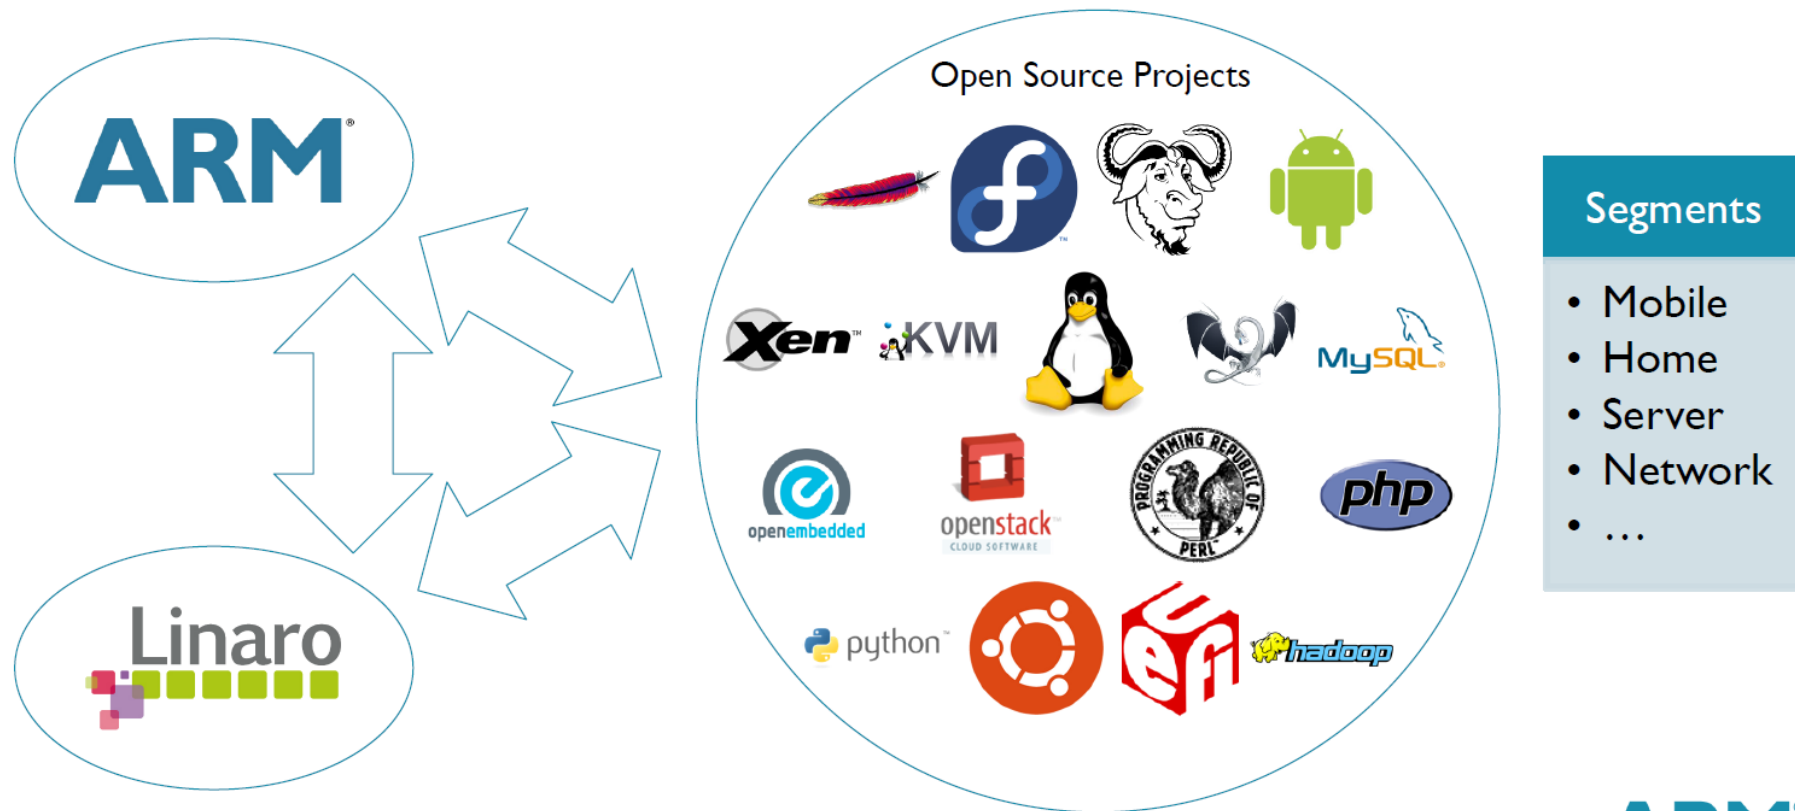
## Firmware & Network Boot



## ARM SoC Framework



# Engaging the Open Source community around ARMv8-A




# Web-scale Workloads – Largely Built On Open Source



# Immediate Target Workloads

## Which ones matter to you?



\* Languages: PHP, Ruby, Java, Python, Ruby, Perl, V8 Javascript, Erlang, Go, Scala etc...